# Joshua Gomes

gomesjoshua@gmail.com | https://www.linkedin.com/in/codewitty | https://github.com/codewitty
San Jose, CA | (408) 891-6172

## EDUCATION

University of California, Irvine - Irvine, CA
Bachelor of Science in Computer Science

## TECHNICAL SKILLS

- **Programming Languages:** Golang, Python, C++, C, JavaScript, HTML5, CSS3, MIPS Assembly
- **Technologies/Databases:** GCP, AWS, Spark, Airflow, SQL, Pandas, Hive, SciKit, PostgreSQL, Flask, REST
- **Tools/Systems:** Unix, Terraform, Terragrunt, Git, Hadoop, Tableau, Pytest, Bitbucket, Jira

## WORK EXPERIENCE

**Software Engineer | Codex Health,** Palo Alto, CA **Sep. 2022–Present**

- Utilized GCP serverless architecture and multi-tenancy to deliver scalable cloud software services to a diverse customer base, achieving a 30% reduction in infrastructure costs and improving system scalability by 50%
- Designed and developed system-facing EHR Integration to track customer cohorts with 10,000+ patients, providing actionable insights for improved patient care, resulting in a 25% improvement in patient outcomes
- Implemented asynchronous task execution using cloud functions, optimizing system performance and reducing response time for users by 40%
- Developed and maintained GCP service integrations, including Cloud Pub/Sub, Cloud Functions, Cloud Run, and Cloud Build, to enhance software application functionality and scalability
- Provisioned GCP resources, including project creation, service account setup, and service configuration, using Terraform and Terragrunt, reducing infrastructure deployment time by 75% and improving system maintainability
- Built highly efficient microservices using cloud functions triggered by HTTP requests or Firestore Database Triggers, improving system resilience and reliability while minimizing cost and maintenance overhead
- Set up and configured CI/CD pipelines using Google Cloud Build, including defining build steps, build triggers, and build logs, to enable automated builds, testing, and deployment of software applications on GCP

**Software Engineer, Big Data | Affinity Solutions,** San Jose, CA **Aug. 2019–Aug. 2022**

- Led a project to develop and deploy a DSAR CCPA compliance framework on the Identity Engine, successfully processing millions of records and ensuring company-wide data privacy compliance
- Migrated the DSAR ETL pipeline to AWS using Hive on EMR, S3, Glue, and RedShift, streamlining data processing and storage and improving performance
- Configured and optimized AWS EMR clusters using Hadoop, Hive, and Spark to process and analyze big data workloads, improving data processing speed and performance by 80%
- Designed and implemented data pipelines using AWS services such as Kinesis, Glue, Lambda, and S3 to extract, transform, and load (ETL) large-scale datasets from various sources, processing over 1 TB of data per day and achieving a 50% reduction in ETL processing time
- Configured and managed Apache Airflow on AWS EC2 instances to automate and orchestrate data workflows, enabling scheduling, monitoring, and alerting of data pipeline tasks, resulting in a 50% reduction in manual labor for pipeline management
- Leveraged open-source libraries and AWS services such as Boto3, Lambda, and S3 to develop an automated sales report delivery system using Python. The solution eliminated 4 hours of weekly manual work and significantly improved client satisfaction by ensuring timely and accurate delivery of reports
- Participated in weekly agile sprints to ensure rapid feature delivery, collaborating closely with cross-functional teams to drive project success

## INTERESTS

- Photography, Badminton, Hiking, Chess, Soccer